



Last updated on **Sep 11, 2026**

Databricks SQL release notes 2024

The following outlines the improvements and updates in Databricks SQL from January through December 2024.

November 21, 2024

SQL warehouse system tables (Public Preview)

- The `system.compute.warehouses` table records when SQL warehouses are created, edited, and deleted. You can use the table to monitor changes to warehouse settings, including the warehouse name, type, size, channel, tags, auto-stop, and autoscaling settings. See [Warehouses system table reference](#).

Data discovery

- The **Lineage** tab in Catalog Explorer has been redesigned with an improved UI for entity filtering.

November 13, 2024

Legacy dashboards:

- Resolved an issue where templated tooltips were not displaying detailed content for dual-axis and multi-field axis charts.

November 6, 2024



Human-readable schedule support for Databricks SQL streaming tables and materialized views

Users can now start, create, and alter schedules for streaming tables and materialized views using human-readable syntax instead of CRON scheduling. See [ALTER MATERIALIZED VIEW](#), [ALTER STREAMING TABLE](#), [CREATE MATERIALIZED VIEW](#), and [CREATE STREAMING TABLE](#).

Streaming tables now support time travel queries

You can now use time travel to query previous table versions based on timestamp or table version (as recorded in the transaction log). You may need to refresh your streaming table before using time travel queries. See [Time travel](#).

Time travel queries are not supported for materialized views.

October 31, 2024

User interface updates

New SQL editor (Public Preview)

- You can now run the active SQL query using the keyboard shortcut `Command` (or `Ctrl`) + `Shift` + `Enter`.
- The parameters input area now shows a scrollbar when the text extends outside of the display window.
- Fixed an issue that prevented the query profile details page from opening fully.
- You can now rename queries by typing the new name into the tab title.
- The **Schedule** button is now disabled for queries that have never been saved before.

October 24, 2024



Release notes for AI/BI tools

The release notes for AI/BI dashboards and Genie Agents (formerly Genie Spaces) have moved to [AI/BI and Genie One release notes](#). Future releases and updates will be documented there.

October 17, 2024

Notification destinations are now generally available

You can create and configure notification destinations that workspace users can add to certain workflows, like alerts, Lakeflow Jobs, and AI/BI dashboard schedules, to send emails or webhooks when an event runs. See [Manage notification destinations](#).

October 10, 2024

Materialized views and streaming tables are now generally available on Databricks SQL

Streaming tables enable incremental ingestion from cloud storage and message queues, while materialized views are pre-computed views that are automatically and incrementally updated as new data arrives. See [Use standalone materialized views](#) and [Use standalone streaming tables](#).

Query insights

- The new columns `query_source`, `executed_as`, and `executed_as_user_id` have been added to the query history system table. See [Query history system table reference](#).

October 3, 2024



User interface updates

The features listed in this section are independent of the SQL warehouse compute versions described above.

Catalog Explorer

- AI-generated comments are now supported for catalogs, schemas, volumes, models, and functions and users can use the inline chat Assistant to help edit their comments.

SQL AI functions The `vector_search()` function is now available in Public Preview. See [vector_search function](#)

September 26, 2024

September 11, 2024

User interface updates

The features listed in this section are independent of the SQL Warehouse compute versions described above.

SQL editor

You can now use named parameter marker syntax in the SQL editor. Named parameter marker syntax can be used across the SQL editor, notebooks, and AI/BI dashboards. See [Use named parameter markers](#).

Queries and legacy dashboards

For SQL queries and legacy dashboards, deleted items no longer appear in the listing pages. Find deleted items in the workspace trash folder. Workspace admins can view deleted items in all users' trash folders.



September 5, 2024

Changes in 2024.40

Databricks SQL version 2024.40 includes the following behavioral changes, new features, and improvements.

Behavioral changes

- **Change to the default schema binding mode for views**

Views now adapt to schema changes in the underlying query by using schema compensation with regular casting rules. This is a change from the previous default of `BINDING` mode, which raised errors when a safe cast could not be performed when referencing the view.

See [CREATE VIEW](#) and `cast` function.

- **Disallow using the undocumented `!` syntax instead of `NOT` outside boolean expressions**

With this release, the use of `!` as a synonym for `NOT` outside of boolean expressions is no longer allowed. For example, statements such as the following: `CREATE ... IF ! EXISTS`, `IS ! NULL`, a `! NULL` column or field property, `! IN` and `! BETWEEN`, must be replaced with: `CREATE ... IF NOT EXISTS`, `IS NOT NULL`, a `NOT NULL` column or field property, `NOT IN` and `NOT BETWEEN`.

This change ensures consistency, aligns with the SQL standard, and makes your SQL more portable.

The [boolean prefix operator](#) `!` (for example, `!is_mgr` or `!(true AND false)`) is unaffected by this change.

- **Disallow undocumented column definition syntax in views**

Databricks supports [CREATE VIEW](#) with named columns and column comments. Previously, the specification of `data type`, `NOT NULL` constraints, or `DEFAULT` has been allowed. With this release, you can no longer use this syntax.

This change ensures consistency, aligns with the SQL standard, and supports future enhancements.

- **Adding a `CHECK` constraint on an invalid column now returns the `UNRESOLVED_COLUMN.WITH_SUGGESTION` error class**

To provide more useful error messaging, in Databricks Runtime 15.3 and above, an `ALTER TABLE ADD CONSTRAINT` statement that includes a `CHECK` constraint referencing an invalid column name returns the `UNRESOLVED_COLUMN.WITH_SUGGESTION` error class. Previously, an `INTERNAL_ERROR` was returned.

New features and improvements

- **Enable UniForm Iceberg using `ALTER TABLE`**

You can now enable UniForm Iceberg on existing tables without rewriting data files. See [On an existing table](#).

- **UTF-8 validation functions**

This release introduces the following functions for validating UTF-8 strings:

- `is_valid_utf8` verified whether a string is a valid UTF-8 string.
 - `make_valid_utf8` converts a potentially invalid UTF-8 string to a valid UTF-8 string using substitution characters
 - `validate_utf8` raises an error if the input is not a valid UTF-8 string.
 - `try_validate_utf8` returns `NULL` if the input is not a valid UTF-8 string.
- **`to_avro` and `from_avro` functions**
- The `to_avro` and `from_avro` functions allow conversion of SQL types to Avro binary data and back.
- **`try_url_decode` function**

This release introduces the `try_url_decode` function, which decodes a URL-encoded string. If the string is not in the correct format, the function returns `NULL` instead of raising an error.

- **Optionally allow the optimizer to rely on unenforced foreign key constraints**

To improve query performance, you can now specify the `RELY` keyword on `FOREIGN KEY` constraints when you `CREATE` or `ALTER` a table.

- **Support for dropping the check constraints table feature**

Selective overwrites using `replaceWhere` now run jobs that delete data and insert new data in parallel, improving query performance and cluster utilization.

- **Parallelized job runs for selective overwrites**

Selective overwrites using `replaceWhere` now run jobs that delete data and insert new data in parallel, improving query performance and cluster utilization.

- **Improved performance for change data feed with selective overwrites**

Selective overwrites using `replaceWhere` on tables with change data feed no longer write separate change data files for inserted data. These operations use a hidden `_change_type` column present in the underlying Parquet data files to record changes without write amplification.

- **Improved query latency for the COPY INTO command**

This release includes a change that improves the query latency for the `COPY INTO` command. This improvement is implemented by making the loading of state by the RocksDB state store asynchronous. With this change, you should see an improvement in start times for queries with large states, such as queries with a large number of already ingested files.



User interface updates

The features listed in this section are independent of the SQL Warehouse compute versions described above.

Visualizations

- Tooltip totals now appear only for stacked charts.

August 22, 2024

Visualizations

For grouped and multi-field configurations, tooltips now show totals when you hover over chart elements.

August 15, 2024

Visualizations

Fixed an issue where row numbers in table visualizations didn't update after changing the page size.

Data discovery

The ability to expand and collapse nested complex column types in Unity Catalog tables is now supported.

August 1, 2024

Visualizations:

- Table sorting is now preserved when data changes due to filtering.

SQL Editor:

- Increased readability by adding additional padding between the last line of a query and the result output.

July 25, 2024

Databricks REST API:

- APIs for managing queries, alerts, data sources, and permissions have changed. The legacy version will continue to be supported for six months. This transition period is intended to give you sufficient time to migrate your applications and integrations to the new version before the older version is phased out. See [Update to the latest Databricks SQL API version](#)

July 18, 2024

User interface updates

Catalog explorer:

- A new catalog configuration wizard is now available for setting up workspace bindings, catalog privileges, and metadata when creating a catalog.

SQL Warehouse monitoring:

- CAN MONITOR permission is now generally available. It allows privileged users to monitor SQL warehouses, including the associated query history and query profiles. See [SQL warehouse ACLs](#).

July 11, 2024



Databricks SQL Version 2024.35 Available

Rollout Schedule: Preview rollout for 2024.35: Between July 9th and July 15th

Changes in 2024.35

Disable column mapping with drop feature

You can now use `DROP FEATURE` to disable column mapping on Delta tables and downgrade the table protocol. See [Maintain compatibility](#).

Variant type syntax and functions in Public Preview

Built-in Apache Spark support for working with semi-structured data as `VARIANT` type is now available in Spark DataFrames and SQL. See [Query variant data](#).

Variant type support for Delta Lake in Public Preview

You can now use `VARIANT` to store semi-structured data in tables backed by Delta Lake. See [Variant type support for Apache Iceberg and Delta Lake](#).

Support for different modes of schema evolution in views

`CREATE VIEW` and `ALTER VIEW` now allow you to set a schema binding mode, enhancing how views handle schema changes in underlying objects. This feature enables views to either tolerate or adapt to schema changes in the underlying objects. It addresses changes in the query schema resulting from modifications to object definitions.

Performance improvement for some window functions

This release includes a change that improves the performance of some Spark window functions, specifically functions that do not include an `ORDER BY` clause or a `window_frame` parameter. In these cases, the system can rewrite the query to run it using an aggregate function. This change allows the query to run faster by using partial aggregation and avoiding the overhead of running window functions. The Spark configuration parameter `spark.databricks.optimizer.replaceWindowsWithAggregates.enabled` controls this optimization and is set to `true` by default. To turn this optimization off, set `spark.databricks.optimizer.replaceWindowsWithAggregates.enabled` to `false`.

Support for the `try_mod` function added

This release adds support for the PySpark `try_mod()` function. This function supports the ANSI SQL-compatible calculation of the integer remainder by dividing two numeric values. If the divisor argument is 0, the `try_mod()` function returns null instead of throwing an exception. You can use the `try_mod()` function instead of `mod` or `%`, which throws an exception if the divisor argument is 0 and ANSI SQL is enabled.

User interface updates

SQL Editor:

The inline assistant is now available in the SQL editor. Click the assistant icon in the editor box to toggle the input. Type a question or comment in English, then press Enter (not Shift+Enter, which runs a query) to generate a response with a different view directly in the editor.

Platform:

An API for notification destinations is now available. You can now programmatically manage webhook and email destinations for your alerts and job run notifications. See [Notification Destinations](#).

July 4, 2024



Databricks SQL Version 2024.30 Available

Rollout schedule: Preview rollout for 2024.30: Between July 9th and July 15th

Changes in 2024.30

Lakehouse Federation is generally available (GA)

[Lakehouse Federation](#) connectors across the following database types are now generally available (GA):

- MySQL
- PostgreSQL
- Amazon Redshift
- Snowflake
- Microsoft SQL Server
- Azure Synapse (SQL Data Warehouse)
- Databricks

This release also introduces the following improvements:

- Support for single sign-on (SSO) authentication in the [Snowflake](#) and [Microsoft SQL Server](#) connectors.
- Stable egress IP support in serverless compute environments. See [Serverless compute firewall configuration](#).
- Support for additional pushdowns (string, math, miscellaneous functions).
- Improved pushdown success rate across different query shapes.
- Additional pushdown debugging capabilities:
 - The `EXPLAIN FORMATTED` output displays the pushed-down query text.

- The query profile UI displays the pushed-down query text, federated node identifiers, and JDBC query execution times (in verbose mode). See [View system-generated federated queries](#).

DESCRIBE HISTORY now shows clustering columns for tables that use liquid clustering

When you run a `DESCRIBE HISTORY` query, the `operationParameters` column shows a `clusterBy` field by default for `CREATE OR REPLACE` and `OPTIMIZE` operations. For a Delta table that uses liquid clustering, the `clusterBy` field is populated with the table's clustering columns. If the table does not use liquid clustering, the field is empty.

Support for primary and foreign keys is generally available

Support for primary and foreign keys in Databricks Runtime is generally available. The GA release includes the following changes to the privileges required to use primary and foreign keys:

- To define a foreign key, you must have the `SELECT` privilege on the table with the primary key that the foreign key refers to. You do not need to own the table with the primary key, which was previously required.
- Dropping a primary key using the `CASCADE` clause does not require privileges on the tables that define foreign keys that reference the primary key. Previously, you needed to own the referencing tables.
- Dropping a table that includes constraints now requires the same privileges as dropping tables that do not include constraints.

To learn how to use primary and foreign keys with tables or views, see [CONSTRAINT clause](#), [ADD CONSTRAINT clause](#), and [DROP CONSTRAINT clause](#).

Liquid clustering is GA

Support for liquid clustering is now generally available using Databricks Runtime 15.2 and above. See [Use liquid clustering for tables](#).

Type widening is in Public Preview

You can now enable type widening on tables backed by Delta Lake. Tables with type widening enabled allow changing the type of columns to a wider data type without rewriting underlying data files. See [Type widening](#).

Schema evolution clause added to SQL merge syntax

You can now add the `WITH SCHEMA EVOLUTION` clause to a SQL merge statement to enable schema evolution for the operation. See [MERGE with schema evolution using SQL, Python, and Scala](#).

Vacuum inventory support

You can now specify an inventory of files to consider when running the `VACUUM` command on a Delta table. See the [OSS Delta docs](#).

Support for Zstandard compression functions

You can now use the `zst_compress`, `zstd_decompress`, and `try_zstd_decompress` functions to compress and decompress `BINARY` data.

Query plans in the SQL UI now correctly display `PhotonWriteStage`

When displayed in the SQL UI, `write` commands in query plans incorrectly showed `PhotonWriteStage` as an operator. With this release, the UI is updated to show `PhotonWriteStage` as a stage. This is a UI change only and does not affect how queries are run.



User interface updates

API support:

- You can now manage notification destinations using the REST API. See [Notification destinations](#).

June 27, 2024

Row Filters and Column Masks in Databricks SQL materialized views and streaming tables are Public Preview

Row filters and column masks in Databricks SQL materialized views and streaming tables are [Public Preview](#). The Public Preview release includes the following changes:

- You can add row filters and column masks to a Databricks SQL materialized view or streaming table.
 - You can define Databricks SQL materialized views or streaming tables on tables that include row filters and column masks.

See [CREATE MATERIALIZED VIEW](#) and [CREATE STREAMING TABLE](#).

User interface updates

Visualizations:

- Improved interactivity in displaying tooltips when hovering over pie, scatter, and heatmap charts with many data points.

Catalog Explorer:



- A revamped Catalog Explorer UI makes it easier to discover and favorite recent Unity Catalog assets from the **Quick Access** view. The navigation experience has also been simplified, allowing you to explore compute, storage, credentials, connections, DBFS, and management details using the  **Settings** in the upper-left corner of the screen. Delta Sharing, Clean Rooms, and External Data now have dedicated pages.

June 6, 2024

Fix for Databricks SQL materialized views and streaming tables

The issue causing `ALTER SCHEDULE` queries on Databricks SQL materialized views and streaming tables to take effect only after the next `REFRESH` operation has been fixed. Now, `ALTER SCHEDULE` queries are applied immediately. See [Schedule refreshes](#).

Materialized views and streaming tables in Databricks SQL are Public Preview

Materialized views and streaming tables in Databricks SQL are Public Preview and available to all customers. The public preview release includes the following changes:

- `REFRESH` of materialized views and streaming tables in Databricks SQL is now synchronous by default. See [REFRESH \(MATERIALIZED VIEW or STREAMING TABLE\)](#).
 - Errors that occur during a refresh operation of a Databricks SQL materialized view or streaming table are returned in the [SQL Editor](#).

To learn how to use materialized views and streaming tables in Databricks SQL, see [Use standalone materialized views](#) and [Use standalone streaming tables](#).

User interface updates



Dashboards:

- Account users can now download visualization data from published dashboards
- Unpublished dashboards can now be published using the **Draft/Publish** dropdown near the top of a dashboard.
- Fixed an issue where parameters named **limit** were not detected.
- Dashboards now appear in the side navigation for AWS GovCloud.

Visualizations:

- Improved box-plot rendering in dark mode.

Query insights:

- For all notebooks attached to SQL warehouses, you can access the query profile by clicking **See performance** under the cell that contains the query. If the cell includes multiple queries, a link to the query profile for each is provided for each statement.

May 30, 2024

New permission level for SQL warehouses

Can monitor permission allows users to monitor SQL warehouses, including the associated query history and query profiles. The **Can monitor** permission is now in [Public Preview](#). See [SQL warehouse ACLs](#).

User interface updates

Dashboard improvements:



- Added a kebab menu to dashboard widgets for downloading images and data.
- Users can migrate legacy dashboard query-backed parameters and dropdown list parameters to dashboards.

May 23, 2024

User interface updates

- You can now select multiple items in the workspace to move or delete. When multiple objects are selected, an action bar appears and provides options to move or delete items. Additionally, you can select multiple items using your mouse and drag them to a new location. Existing permissions on objects still apply during bulk move and delete operations.
- You can now mark Unity Catalog assets as favorites in the Catalog Explorer and Schema Browser. This includes catalogs, schemas, tables, models, volumes, and functions. Unity Catalog assets that you mark as favorites are easily accessible from the Databricks homepage.

Dashboard updates:

- Dual-axis combo charts now correctly display bar legends on the right axis and line legends accordingly.
- Dual-axis charts now correctly show labels on bars.

Visualizations updates:

- The table editor's conditional format labels for **if** and **then** now support dark mode.
- The redundant **open link** icon has been removed from the table editor's format tooltips.
- The default font color's label in the table editor now aligns automatically.

May 16, 2024

Rollout schedule

- Preview rollout for 2024.25: Completed May 1st
- Current rollout for 2024.25: Between May 14th and May 21st

NOTE

An upgrade to the panda Python library (version 2.0.3) caused a breaking change in Databricks SQL version 2024.20. Databricks did not roll out version 2024.20 to the current channel. Instead, the preview channel was upgraded to 2024.25 on May 1, 2025. The current channel rollout goes directly from 2024.15 to 2024.25.

Changes in 2024.25

Data governance

Credential passthrough and Hive metastore table access controls are deprecated.

[Credential passthrough](#) and [Hive metastore table access controls](#) are legacy data governance models. Upgrade to Unity Catalog to simplify the security and governance of your data by providing a central place to administer and audit data access across multiple workspaces in your account. See [What is Unity Catalog?](#).

Support for credential passthrough and Hive metastore table access controls will be removed in an upcoming Databricks Runtime version.

SQL language features

The * (star) clause is now supported in the WHERE clause.

You can now use the star (*) clause in the `WHERE` clause to reference all columns from the `SELECT` list. For example, `SELECT * FROM VALUES(1, 2) AS T(a1, a2) WHERE 1 IN(T.*)`.

Support for Cloudflare R2 storage to avoid cross-region egress fees since 2024.15

You can now use Cloudflare R2 as cloud storage for data registered in Unity Catalog. Cloudflare R2 is intended primarily for Delta Sharing use cases in which you want to avoid the data egress fees charged by cloud providers when data crosses regions.

Cloudflare R2 storage supports all of the Databricks data and AI assets supported in AWS S3.

See [Use Cloudflare R2 replicas or migrate storage to R2](#) and [Connect to a Cloudflare R2 external location](#).

User interface updates

The features listed in this section are independent of the SQL Warehouse compute versions described above.

Data discovery updates: The Hive metastore to Unity Catalog update wizard supports upgrading Hive metastore managed tables using all-purpose compute or SQL warehouses. Updating more than 20 tables creates a new notebook that contains the `SYNC` AND `ALTER TABLE` commands that perform the conversion.

Dashboard updates:

- Dual-axis functionality is now available for Area, Bar, Line, and Scatter chart types.
 - When you enable a dual-axis chart, the axis title and range is no longer copied to the secondary axis.
 - The last field identified in the visualization configuration is automatically relocated to the right-side y-axis.

May 9, 2024

SQL Editor fixes:



- The admin setting **Results table clipboard features** now applies to the SQL editor's **New result table**.

Dashboard improvements:

- Query-based parameters allow authors to define a list of selectable values that viewers can use as parameters for other visualizations on a dashboard canvas. See [Use query-based parameters](#).
- Column order in files downloaded from a table widget is now preserved.
- The table editor now includes hover tooltips that display the names of columns.
- When switching from other visualization types to a histogram, information encoding is now better preserved.

Dashboard fixes:

- Fixed an issue where a single grid height filter displayed an unnecessary overflow scrollbar.
- Fixed an issue that caused incorrectly rendered visualizations on published dashboards where a referenced dataset column was deleted.

May 2, 2024

Serverless SQL Warehouse support expanded: Serverless SQL warehouses are now available in the following regions:

- ca-central-1
- ap-northeast-2

See [Features with limited regional availability](#)

Dashboard updates:

- Queries and visualizations can now be copied to a new dashboard from SQL editor. You can still add visualizations to legacy dashboards from the SQL editor. See [Edit, download, or add to a dashboard](#).

- Dashboards will now maintain a 24-hour result cache to optimize initial loading times. See [Dataset optimization and caching](#).
- Bar charts with categorical X and quantitative Y are now sorted in Databricks Assistant responses.
- Corrected migration issue with legacy histogram `COUNT (*)` to ensure accurate migration.
- Implemented the ability to mix numeric types and date types in a filter widget.
- When creating charts, Databricks Assistant now automatically suggests relevant columns as you type.

Visualization updates:

- User-selected color for tables now persists across light and dark modes in legacy charts.
- Data truncation logic has been improved to enhance performance in combo, pie, heatmap, and histogram charts.
- A tick mark is now always displayed at the top of a quantitative axis for basic charts.

April 23, 2024

UI updates:

- For all **Share** dialogs in the UI, the `All Users` group has been renamed to `All Workspace Users`. The new name more accurately reflects the scope of the group, which has always included users assigned to the workspace. No change is made to group membership as part of this rename.

Dashboard improvements:

- When a dashboard's SQL warehouse is starting, a dialog appears to explain the wait time.
- Scroll position is preserved when switching between the **Canvas** and **Data** tabs.
- Cloning a legacy dashboard to create a Lakeview dashboard now supports some parameter conversion. See [Adjust legacy parameters](#).



- Relative dates, such as **Today**, **Tomorrow**, and **7 days ago**, are now supported for date and date time parameters.
- Number range sliders can be added as filters on a dashboard.
- Histograms can now display disaggregated data.
- Scatter plots now support size encoding.

Dashboard fixes:

- Temporal color encoding can now change color assignments correctly.

Visualization updates:

- Custom tooltip formats now function correctly for multi-axis charts.
- The **New charts** preview tag is removed when users have not changed the toggle in the past 14 days.

April 18, 2024

Lakeview dashboards are generally available

Lakeview dashboards are now the default dashboarding tool. They have been renamed as **Dashboards** in the UI. Databricks SQL dashboards are now called **Legacy dashboards**. The names of the related API tools have not changed.

- **Dashboard improvements:**

- Audit logs are available for Lakeview dashboards. See [AI/BI dashboard events](#).
- Data downloaded from dashboards respect applied parameters.
- Databricks Assistant is enabled on the **Data** tab without adding datasets first.
- Stacked bar charts with multiple Y fields can support sorting the X-axis based on the sum of Y-axis values.
- Toggle between **Linear** and **Log(Symmetric)** scale functions in visualization axis menus.



- The default size of the filter widget is now more compact.
- The initial load time for the text-entry filter widget has been reduced.
- Improved automatic chart conversions when migrating from legacy dashboards.
- **Dashboard fixes:**
 - The restricted viewing settings warning does not show if the dashboard has been shared with others.
 - The error messages in the **Data** tab SQL editor are now dark mode enabled.

User interface updates

The features listed in this section are independent of the SQL warehouse compute versions described above.

Improvements:

- The tooltips on stacked charts now display the stack value and percentage by default.
- The tooltips for multi-axis charts now highlight the hovered item.
- **Table** visualizations for Databricks SQL now adapt a new query result's data type when edited in the SQL editor.
- The Catalog Explorer's **Query History** table shows a tree-like view for **Query Source** attribution. You can use this to see which entities have triggered the query statement to run.

April 11, 2024

User interface updates

The features listed in this section are independent of the SQL warehouse compute versions described above.

Improvements:



- You can now group by percentage when creating visualizations in Databricks SQL and notebooks.
- For new charts (in Public Preview), you can zoom in along a single axis by clicking and dragging in a straight line parallel to the axis.
- The Unity Catalog shared cluster **Allowlist** UI is now generally available. You can access it on the Metastore details page in Catalog Explorer. See [How to add items to the allowlist](#).
- Forms to create and edit external locations now open as a full page. They include the option to include a storage credential.

Fixes:

- Corrected an issue for Histogram charts where negative values were erroneously marked as positive.

April 4, 2024

User interface updates

The features listed in this section are independent of the SQL warehouse compute versions described above.

Improvements:

- Improvements to Histogram charts on Lakeview dashboards.
 - Added support for labels.
 - Bin settings are now retained when switching between different fields.
- The samples gallery on the dashboard listing page now creates Lakeview dashboards. See [Tutorial: Use sample dashboards](#).
- Right-clicking on the border of a widget on a Lakeview dashboard opens a context menu.
- The left-side navigation bar is retained for workspace users viewing published Lakeview dashboards.
- Filter selections are retained when navigating between published and draft Lakeview dashboards.

- Column names can now be inserted into the SQL editor when editing a query from the **Data** tab in a draft Lakeview dashboard.
- Replacing a Lakeview dashboard keeps the existing dashboard name and replaces the contents.
- Switching visualizations between heat maps and other chart types now preserves the relevant fields better.

Fixes:

- Bar charts with color encodings now correctly restrict adding multiple Y-axis fields.
- Resolved an issue where the **Download as PNG** button was missing from some visualizations.
- Corrected formatting for negative big integers previously missing thousands of separators.
- Fixed incorrect hover line placement when hovering over labels on line charts.

March 28, 2024

Databricks SQL version 2024.15 available

Rollout Schedule

- Preview rollout for 2024.15: Between February 14th and March 4th
- Current rollout for 2024.15: Between March 25th and April 8th

Changes in 2024.15

Delta updates

- **Delta UniForm is now generally available:** UniForm is now generally available and uses the IcebergCompatV2 table feature. You can now enable or upgrade UniForm on existing tables. See [Read Delta Lake tables with Iceberg clients](#).



- **Recompute data skipping statistics for Delta tables:** You can now recompute statistics stored in the Delta log after changing columns used for data skipping. See [Specify statistics columns](#).

SQL language updates

- **Declare temporary variables in a SQL session:** This release introduces the ability to declare temporary variables in a session that can be set and then referred to from in queries. See [Variables](#).
- **Native XML file format support (Public Preview):** Native XML file format support is now in Public Preview. XML file format support enables ingestion, querying, and parsing of XML data for batch processing or streaming. It can automatically infer and evolve schema and data types, supports SQL expressions like `from_xml`, and can generate XML documents. It doesn't require external jars and works seamlessly with Auto Loader, `read_files`, `COPY INTO`, and DLT. See [Read and write XML files](#).
- **Cloud Fetch is now enabled by default:** Cloud Fetch is enabled by default in AWS workspaces with bucket versioning enabled. If you have bucket versioning enabled, Databricks recommends setting a lifecycle policy to remove old versions of uploaded query results automatically. See [Cloud Fetch in ODBC \(ODBC\)](#) and [Cloud Fetch in JDBC \(JDBC\)](#).

Apache Spark SQL updates

Databricks SQL 2024.15 include Apache Spark 3.5.0. Additional bug fixes and improvements for SQL are listed on the Databricks Runtime 14.3 release note. See [Apache Spark](#) and look for the `[SQL]` tag for a complete list.

User interface updates

The features listed in this section are independent of the SQL warehouse compute versions described above.

- A new overview tab in the entity page of Catalog Explorer shows important metadata like filesize, data source, owner, table schema, and comments.
- **Lakeview dashboard updates:**



- Lakeview dashboards now support parameters. Authors can add parameters to dataset queries from the **Data** tab and then set parameters on the canvas using single-value selectors and date-pickers. See [Work with dashboard parameters](#).
- Lakeview dashboards are now supported in the workspace permissions API. See [PATCH /api/workspace/workspace/updatepermissions](#) in the REST API reference.
- **Control** widgets on the canvas have been renamed to **Filter** widgets.
- Combo charts no longer allow disaggregated fields on the x-axis.
- The **Copy link** button in the Share dialog now includes parameters stored in the URL.
- Widgets on published dashboards no longer show borders when hovering.
- Resolved an issue where the Databricks Assistant and Download PNG buttons overlap with chart visuals.
- Switching visualizations between heatmaps and other chart types now better preserves the relevant fields.
- Bar charts with color encodings now correctly restrict adding multiple Y-axis fields.

March 21, 2024

- Unity Catalog model lineage is now in Public Preview.

The table view in Catalog Explorer now has an **Overview** tab to describe its primary metadata.

- SQL warehouses for notebooks, now generally available, allow you to take advantage of fully managed, instant, and scalable compute for your SQL workloads in the rich, collaborative authoring environment of a notebook. See [Use a notebook with a SQL warehouse](#).

The following fixes and improvements apply to Lakeview dashboards:

- Expanded API support for Lakeview adds the ability to create, get, update, and trash dashboards. See [Lakeview](#) in the REST API reference.



- Added a refresh button for the **Catalog** browser on the **Data** tab.
- **Lakeview dashboards** now appear before **Dashboards** in the **New** menu in the workspace sidebar. On the dashboard listing page, the **Lakeview dashboards** tab appears to the left of the **Dashboards** tab.
- The Databricks Assistant experience for Lakeview has been updated with an input box and suggestions to improve discoverability and help users understand the prompts they can pose.
- Lakeview visualizations now support median aggregations.
- Updated the color picker in the Lakeview dashboard visualization editor for a more streamlined user experience when creating tables.
- Improved pie chart migration to exclusively support scenarios with angle or color definitions.
- Fixed a bug preventing grouping by charts named **count**. Charts can now be grouped by fields named **count**.
- For bar charts, **group** and **stack** layout controls are now hidden when not applicable to the user-selected configuration.

March 14, 2024

For Lakeview dashboards:

- Histograms now support custom categorical colors.
 - Heatmaps now support quantitative scales.
 - Titles and descriptions are retained when switching between visualization types, including Combo charts.
 - You can now open the underlying dataset associated with a draft dashboard widget by right-clicking on it. The dataset opens in the **Data** tab.
- New charts now apply aliases and custom colors for null values in numeric columns.
- New charts now render tick marks to show the top of the y-axis.

March 7, 2024

- When viewing a table in Catalog Explorer, the **Create** button includes an option to create a Lakeview dashboard rather than a Databricks SQL dashboard.
- Histograms are now available for Lakeview dashboards. Histograms are commonly used to visualize the distribution of a numeric field.
- When cloning a Databricks SQL dashboard to create a Lakeview dashboard, dataset conversion issues now show as errors in the new widget on the Lakeview dashboard.
- Color gradients are available when a numerical field is used for a visualization on a Lakeview dashboard.
- Color gradients are now exposed in the Lakeview dashboard visualization editor when a **Color by field** is specified.
- The title and description associated with a visualization no longer appear editable if the viewer lacks editing privileges on a draft Lakeview dashboard.
- Fixed an issue where tooltips in charts with over 100 series incorrectly showed all series. Now, only the focused series is shown.
- Reduced typing latency in the SQL editor by 30% through performance optimizations.
- When managing queries in the SQL editor, moving a query to trash automatically closes the tab.
- Fixed an issue in the SQL editor where text was accidentally selected when adjusting the side panel width.

February 29, 2024

- Serverless SQL warehouse support has been added in the following regions:
 - ap-south-1
 - ap-southeast-1
 - ap-northeast-1
 - sa-east-1
 - eu-west-3

See [Databricks clouds and regions](#) for a complete list of supported regions.

- The schema browser in Catalog Explorer now displays column primary and foreign key constraints.
- The retention time shown in the Lineage tab in Catalog Explorer has been increased to one year.
- Tooltips on new charts in notebooks are now always rendered inside the visualization boundary.
- Learn how to programmatically manage Lakeview dashboards using the REST API. See [Manage dashboards with Workspace APIs](#).
- Lakeview dashboards now support histograms.
- Improved sharing and publishing in Lakeview dashboards:
 - Improved share and publish dialogs, allowing safe and easy sharing to any account user.
 - Dashboards opened from the workspace browser show the published dashboard if it exists. Viewers can now also see details of the latest published version, including publisher, time, and credentials.
 - For editors, a new drop-down switcher in the Lakeview Dashboard UI allows you to quickly move between draft and published versions.

February 22, 2024

- Improvements to the **Sample Data** tab in the Catalog Explorer table view enable you to sort columns, **Copy** selected data to your clipboard, and view line numbers. It can now better display special values, like JSON objects, dates, numeric and null values.
- Lakeview dashboards now support sending periodic PDF snapshots of the dashboard to workspace users and notification destinations. See [Manage scheduled dashboard updates and subscriptions](#).
- The list of visualization options in the Lakeview dropdown picker is now sorted alphabetically.
- When copying Databricks SQL dashboards to Lakeview dashboards, widgets that cannot be converted now show the visualization configuration picker instead of an error message.

February 15, 2024

- The documentation for code-based query filters, such as `SELECT action AS 'action::filter'`, has been removed. Databricks recommends updating queries to remove this pattern.
- For Lakeview dashboards, pie charts now display equal-sized slices when no angle field is specified.
- Lakeview now supports combo charts, which combine bar and line charts to show two different values on the same chart.
- Heatmap charts, which use color intensity to show the magnitude of the correlation between two discrete variables, are now available in Lakeview.

February 8, 2024

- You can now request access when opening a link to a Lakeview dashboard you do not have permissions on.
- Lakeview dashboard filters now have explicit **All** and **None** options. Authors can choose to hide the **All** option in single select filters.
- You can now set minimum and maximum values for axes on Lakeview dashboard charts.

February 1, 2024

Databricks SQL Version 2024.10 Available

Rollout Schedule

- Preview rollout for 2024.10: Between Jan 30, 2024 and Feb 5, 2024
- Current rollout for 2024.10: Between Feb 13, 2023 and Feb 20, 2024

Changes in 2024.10



- **Fixed corrupt file handling in DML commands:** The DML commands `DELETE`, `UPDATE`, and `MERGE INTO` no longer respect the read options `ignoreCorruptFiles` and `ignoreMissingFiles`. When encountering an unreadable file in a table, these commands now fail even if these options are specified.
- **Row-level concurrency is Generally Available and on by default:** Row-level concurrency reduces conflicts between concurrent write operations by detecting changes at the row-level. Row-level concurrency is only supported on tables without partitioning, which includes tables with liquid clustering. Row-level concurrency is enabled by default on Delta tables with deletion vectors enabled. See [Row-level concurrency](#).
- **Shallow clone for Unity Catalog external tables (Public Preview):** You can now use shallow clone with Unity Catalog external tables. See [Shallow clone for Unity Catalog tables](#).
- **Faster multi-threaded statistics collection:** Statistics collection is up to 10 times faster on small clusters when running `CONVERT TO DELTA` or cloning from Iceberg and Parquet tables. See [Convert to Delta Lake](#) and [Incrementally clone Parquet and Apache Iceberg tables to Delta Lake](#).
- **Pushdown filters in the DeltaSource on Delta files:** For better utilization, partition filters on Delta tables streaming queries are now pushed down to Delta before rate limiting.

User interface updates

The features listed in this section are independent of the SQL Warehouse compute versions described above.

- The **Admin view** tab on listing pages for Databricks SQL objects (queries, dashboards, and alerts) has been removed. Workspace admin users can view all objects from their respective listing pages. See [Access and manage saved queries](#), [Legacy dashboards](#), and [Databricks SQL alerts](#).
- The query history page displays queries from the past 24 hours by default. See [Query history](#).
- A menu option, **Clone to Lakeview dashboard**, has been added to the Databricks SQL dashboard UI. You can use this tool to create a new Lakeview dashboard that includes the same queries and visualizations in your existing Databricks SQL dashboards. See [Clone a legacy dashboard to an AI/BI dashboard](#).
- Bar charts in Lakeview dashboards support stacking bars to normalize to 100%.

- Fixed a problem where zooming in on a published Lakeview dashboard resulted in focusing on incorrect zoom intervals.

January 24, 2024

- The Lakeview dashboard canvas automatically adjusts widget placement to remove empty vertical white space between rows when possible.
- Reduced whitespace between title and description text in Lakeview dashboard visualizations.

January 18, 2024

- Fixed a rendering issue for visualizations where bar charts showing a single date on the x-axis resulted in a very thin bar. New chart visualizations render as expected.
- The Lakeview dashboard listing page shows your dashboards by default. You can use filters on that page to access Lakeview dashboards owned by other workspace users.

January 11, 2024

- Databricks SQL Queries and Dashboard APIs support changing the **Run as** role setting programmatically.
- Lakeview supports exporting and importing dashboards as files to facilitate reproducing draft dashboards across workspaces. See [Export, import, or replace a dashboard](#)

January 4, 2024

- Introduced primary key and foreign key entity relationship diagrams in Catalog Explorer. See [View the Entity Relationship Diagram](#).

On this page

 Ask Genie

Was this page helpful?



Yes



No



Send feedback